



MASTER OF
MANAGEMENT
ANALYTICS

MMA 867

Predictive Modelling

Anton Ovchinnikov

Assignment 1 - Section 2

May 3, 2020 11:59 PM

Shangeri Sivalingam

Order of files:

Filename	Pages	Comments and/or Instructions

Additional Comments:

Kaggle name: Bike Sharing Demand

Total number of teams on the leaderboard: 3242

Position on the leaderboard at the time of your last submission: 1446

3 Competition Choices

Option 1: House Prices: Advanced Regression Techniques

Predict sales prices and practice feature engineering, RFs, and gradient boosting. With 79 explanatory variables describing (almost) every aspect of residential homes in Ames, Iowa, this competition challenges you to predict the final price of each home.

Number of entrants: 4751

Status: Ongoing

Link to Kaggle competition: <https://www.kaggle.com/c/house-prices-advanced-regression-techniques>

Option 2: New York City Taxi Trip Duration

Predicts the total ride duration of taxi trips in New York City. Your primary dataset is one released by the NYC Taxi and Limousine Commission, which includes pickup time, geo-coordinates, number of passengers, and several other variables.

Number of entrants: 1254

Status: Completed 3 years ago

Link to Kaggle competition: <https://www.kaggle.com/c/nyc-taxi-trip-duration/data>

Option 3: Bike Sharing Demand

Participants are asked to combine historical usage patterns with weather data in order to forecast bike rental demand in the Capital Bikeshare program in Washington, D.C.

Number of entrants: 3242

Status: Completed 5 years ago

Link to Kaggle competition: <https://www.kaggle.com/c/bike-sharing-demand/data>

Final choice chosen:

Option 3 was chosen - Bike Sharing Demand because it required feature engineering and advanced regression techniques. This competition provided good practice and exposure for predictive modelling, for example, the usage of Extreme Gradient Boosting. This competition requires a predictive regression model to predict the number of bike rental for a given date, season, holiday, working day, weather, temperature, "feels like" temperature, humidity and wind speed. The following section will be a detailed description of how the predictive model was built. Root Mean Squared Log Error (RMSLE) will be used to evaluate the actual versus predicted.

Building the regression model

Understanding the data

To begin, the data was first checked for missing data (there was no missing data) and the datetime column in the train and test set was spilt into hour, day and month columns to see if these variables explain the number of bike rentals better.

Before building the regression model, a simple graph was plotted to understand any trends in the data. It could be possible that season affects the number of bike rentals. To be sure, the number of bikes was plotted against each hour of the day for each season as seen in Figure 2.1 below.

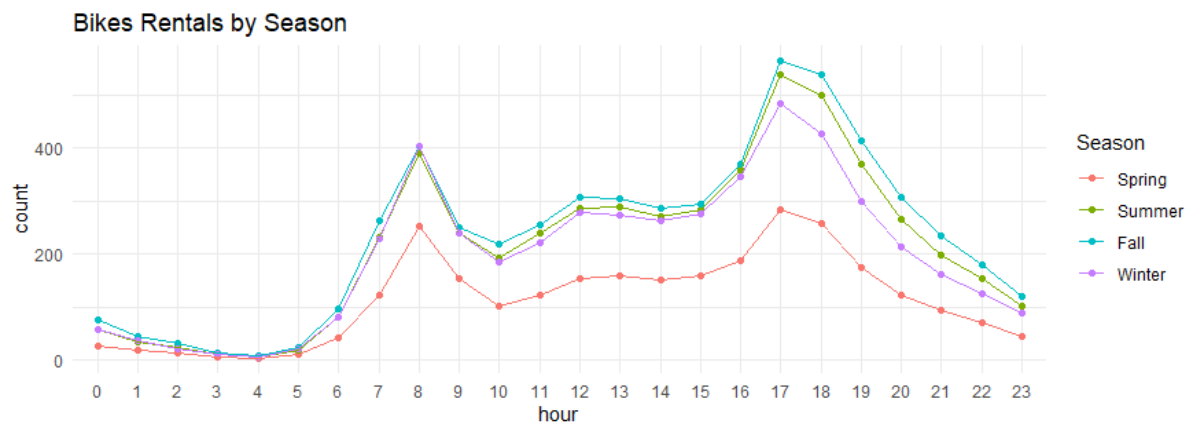


Figure2.1: Line graph of bike rentals against hour by season.

From the graph, we understand there are more bike rental in morning, from the 7th hour and in the evening from the 17th to 18th hour. Furthermore, people rent bikes more in fall, summer and winter, and much less in spring. We will bear variable season and hour in mind when building the model later.

Building the model

Target RMSLE: 0.49644 (position 1625)

The above RMSLE score is the average score on Kaggle for this competition.

Step 1: separating train dataset to 70% for model building and 30% for testing

The following is a breakdown of the variables and their description:

- datetime - hourly date + timestamp
- season - 1 = spring, 2 = summer, 3 = fall, 4 = winter
- holiday - whether the day is considered a holiday
- workingday - whether the day is neither a weekend nor holiday
- weather
 - 1: Clear, Few clouds, Partly cloudy, Partly cloudy
 - 2: Mist + Cloudy, Mist + Broken clouds, Mist + Few clouds, Mist
 - 3: Light Snow, Light Rain + Thunderstorm + Scattered clouds, Light Rain + Scattered clouds
 - 4: Heavy Rain + Ice Pallets + Thunderstorm + Mist, Snow + Fog
- temp - temperature in Celsius
- atemp - "feels like" temperature in Celsius
- humidity - relative humidity
- windspeed - wind speed
- **count - number of total rentals (this is the dependent variable to be predicted)**

Step 2: Start modelling with all factors

Model: `fit<-lm(count~., train_data)`

RMSE: 108.6011

The RMSE score seems to be high. In the case of RMSE, the presence of outliers can explode the error term to a very high value. Although, it is better to over specify than under specify the model, outliers may influence bias. As much as possible, all data points should be accounted for in the

model if possible. There could be a possibility of outliers in this dataset, however, that can be revisited later on if a good enough RMSLE score cannot be achieved.

Step 3: Log the dependent variable

```
Model: logfit<-lm(log(count)~., train_data)
```

RMSLE: 0.613963

Although an improvement, the RMSLE score is not good enough.

Step 4: Using Step AIC to get a better model

```
Model: fit<-lm(log(count)~., train_data)
```

```
logfitAIC <- stepAIC(fit, direction = 'both')
```

RMSLE: 0.6140443

The stepAIC() function from the MASS package to get a better model. The stepAIC() function performs model selection by starting from a "maximal" model, which is then trimmed down to a model with the independent variables that best explain the dependent variable. However, based on the RMSLE, the "improved" model is worse than log(count) ~. in Step 3. Continue with the model in step 3.

Step 5: Using train() from Caret package

```
Model: logfit2 <- train(log(count)~., train_data, method="xgbLinear", trControl = ctrl)
```

RMSLE: 0.422968

Kaggle Score: 0.52418 [position 1902]

To aid in the predictive modelling process, the train function in the Caret package was used. The method xgbLinear of XGBoost was used to build a new regression model. Gradient boosting is an approach where instead of training the models in isolation of one another like in the previous steps, each new model is trained to predict and correct the residuals (errors) made by the previous model. The models are then added together to make the final prediction. Furthermore, as tuning parameters were added to get a better average error term. All in all, a good RMSLE score was achieved so far.

So this new model was used to predict the test dataset to be submitted on Kaggle. However, the score is not good enough to reach the target RMSLE: 0.49644 (position 1625).

Step 6: Remove the outliers

To improve the model, the train dataset was plotted to see if there are any influential outliers to be removed. As shown from the plot below (Figure 2.2), there are some outliers that can potentially be removed to build a better model. These outliers were removed.

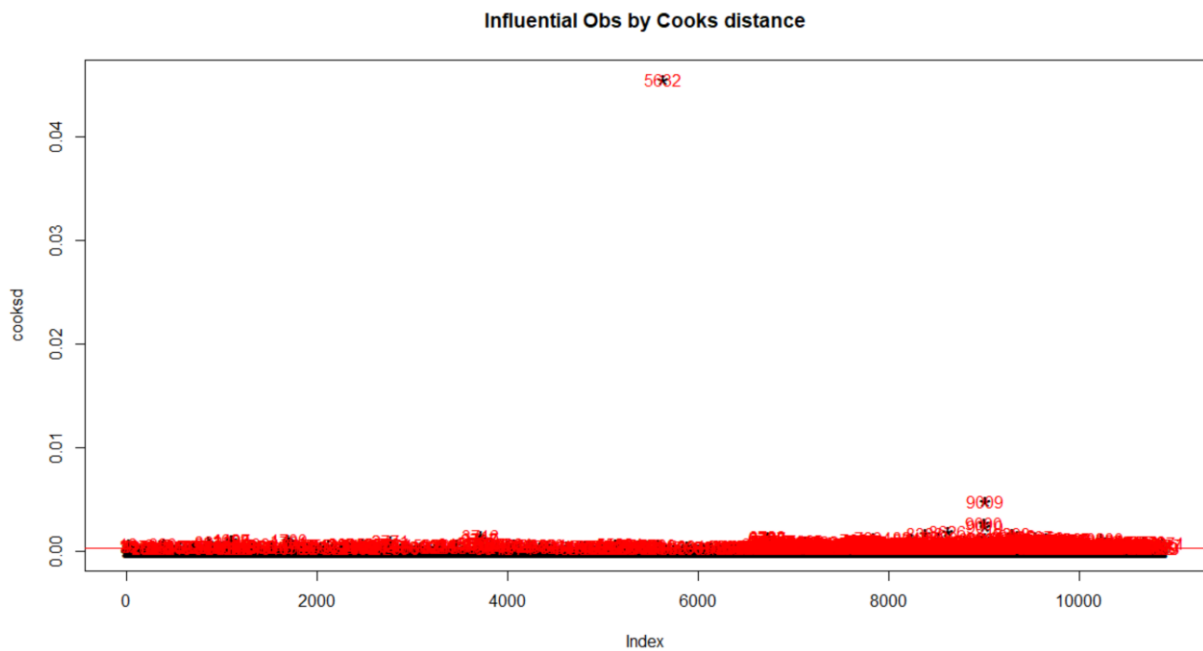


Figure 2.2: Plot of outliers in the train dataset

Step 7: Using XGBoost (xgblinear) on train dataset with no outliers

Model: `logfit2 <- train(log(count)~., train_data, method="xgbLinear", trControl = ctrl)`

RMSLE: 0.4237904

Kaggle Score: 0.50459

The RMSLE and Kaggle score improvement significantly but not good enough.

Step 8: Using `stepAIC()` to find a better model

`fit<-lm(log(count)~., train_data)`

`logfitAIC <- stepAIC(fit, direction = 'both')`

The new model chosen from the `stepAIC()` function is as follows:

`log(count) ~ season + holiday + workingday + weather + temp + humidity + windspeed + hour + day + month`

Step 9: Training the new model [Caret package]

Model: `logfit2 <- train(log(count) ~ season + holiday + workingday + weather + temp + humidity + windspeed + hour + day + month, train_data, method="xgbLinear", trControl = ctrl)`

RMSLE: 0.4231934

While the RMSLE improved, it is not good enough.

Step 10: Log the season and hour variable

Model: `logfit2 <- train(log(count) ~ log(as.numeric(season)) + holiday + workingday + weather + temp + humidity + windspeed + log(as.numeric(hour)) + day + month, train_data, method="xgbLinear", trControl = ctrl)`

RMSLE: 0.3432964

Kaggle Score: 0.49247 [position 1446 on public and private leaderboard]

Target score achieved with this model. Since the competition is completed, the score does not show up on the private or public leaderboard, but nonetheless, the “would be” position on the Kaggle

public and private leaderboard is 1446 (shown in Figure 2.3) which is better than the average target score – 0.49644.

Appendix

Final Kaggle Score

The screenshot shows the Kaggle submission interface. At the top, there's a search bar and navigation tabs: Overview, Data, Notebooks, Discussion, Leaderboard (selected), Rules, Team, My Submissions, and Late Submission. Below the tabs, it says "Your most recent submission" and displays a table with the following data:

Name	Submitted	Wait time	Execution time	Score
predicted_rentals.csv	2 minutes ago	0 seconds	0 seconds	0.49247

Below the table, there's a green bar with the word "Complete" and a link "Jump to your position on the leaderboard".

Below this, there's a section for the "Public Leaderboard" and "Private Leaderboard". The private leaderboard is calculated over the same rows as the public leaderboard in this competition. This competition has completed. This leaderboard reflects the final standings. There's a "Refresh" button.

The public leaderboard table shows the following data:

#	Δpub	Team Name	Notebook	Team Members	Score	Entries	Last
1	—	Bolaka Mukherjee			0.33756	26	5y
2	—	Logical Guess			0.34821	25	5y

Appendix 1: Final Kaggle score: 0.49247

The screenshot shows the Kaggle public leaderboard. At the top, there's a search bar and navigation tabs: Overview, Data, Notebooks, Discussion, Leaderboard (selected), Rules, Team, My Submissions, and Late Submission. Below the tabs, it says "Your most recent submission" and displays a table with the following data:

#	Δpub	Team Name	Notebook	Team Members	Score	Entries	Last
1443	—	Let's Battle			0.49210	14	5y
1444	—	SPSS			0.49212	1	5y
1445	—	Julaff			0.49238	9	6y
1446	—	Victor Rebli			0.49249	7	5y
1447	—	oneday			0.49258	6	5y
1448	—	Juho			0.49272	4	6y
1449	—	Honglin Yu			0.49286	6	5y
1450	—	Awesome Ninjas			0.49293	8	5y
1451	—	aboev			0.49294	15	6y
1452	—	Nitin Reddy			0.49296	4	5y
1453	—	h.k.n			0.49299	4	5y
1454	—	lemur			0.49300	5	5y
1455	—	Pablo			0.49303	5	5y

Appendix 2: Kaggle public leaderboard with "would be" positions highlighted

R Script

```
1 library("readxl")
2 library("lubridate")
3 library("proto")
4 library("gsubfn")
5 library("RSQLite")
6 library("sqldf")
7 library("tidyverse")
8 library("ggplot2")
9 library("mice")
10 library("caret")
11 library("dplyr")
12 library("tidyr")
13 library("MASS")
14 library("car")
15 library("Metrics")
16 library("glmnet")
17 library("xgboost")
18 library("gbm")
19 library("mboost")
20
21 train <- read.csv("./Assignment//Individual/bike-sharing-demand//train.csv")
22 test  <- read.csv("./Assignment//Individual/bike-sharing-demand//test.csv")
23 #Duplicate the test data to save the datetime for later
24 to.predict <- test
25
26 #step 1: lets look at the variables we are dealing with
27 head(train)
28 head(test)
29
30 #columns "casual" and "registered" is not required in the model, so let's remove that from the train
31 head(test)
32
33 #columns "casual" and "registered" is not required in the model, so let's remove that from the train
34 train = train[!(names(train) %in% c("casual"))]
35 train = train[!(names(train) %in% c("registered"))]
36
37 #Step 2: check for missing values
38 md.pattern(train)
39 md.pattern(test)
40 #no missing data!
41
42 #Step 3: Converting integer to factor
43 # #training set
44 train$season <- as.factor(train$season)
45 train$holiday <- as.factor(train$holiday)
46 train$workingday <- as.factor(train$workingday)
47 train$weather <- as.factor(train$weather)
48
49 #test set
50 test$season <- as.factor(test$season)
51 test$holiday <- as.factor(test$holiday)
52 test$workingday <- as.factor(test$workingday)
53 test$weather <- as.factor(test$weather)
54 #
55 #Step 4: let's work with the train set
56 #Deriving day, hour from datetime field
57 train$datetime <- ymd_hms(train$datetime)
58 train$hour <- hour(train$date)
59 train$day <- wday(train$date)
60 train$month <- month(train$date, label=T)
```

```

58
59 #Deriving day, hour from datetime field
60 test$datetime <- ymd_hms(test$datetime)
61 test$hour <- hour(test$date)
62 test$day <- wday(test$date)
63 test$month <- month(test$date, label=T)
64
65 str(train)
66 names(train)
67
68 str(test)
69 names(test)
70
71 train[,11:13]<-lapply(train[,11:13], factor) #converting derived variables into factors
72 test[,10:12]<-lapply(test[,10:12], factor) #converting derived variables into factors
73
74 #Step 4: Removing datetime field
75 train$datetime <- NULL
76 colnames(train)
77 str(train)
78
79
80 #Removing the data field for test
81 test$datetime <- NULL
82
83 #Step 5: visualization so we can understand the data
84 season_summary_by_hour <- sqldf('select season, hour, avg(count) as count from train group by season, hour')
85
86 #There are more rental in morning(from 7 hour) and evening(17-18th hour)
87 #People rent bikes more in Fall Summer and Winter, and much less in Spring

```

```

86 #There are more rental in morning(from 7 hour) and evening(17-18th hour)
87 #People rent bikes more in Fall Summer and Winter, and much less in Spring
88 plot <- ggplot(train, aes(x=hour, y=count, color=season))+
89   geom_point(data = season_summary_by_hour, aes(group = season))+
90   geom_line(data = season_summary_by_hour, aes(group = season))+
91   ggtitle("Bikes Rentals by Season")+ theme_minimal()+
92   scale_colour_hue('Season',breaks = levels(train$season),
93     labels=c('Spring', 'Summer', 'Fall', 'Winter'))
94 plot
95
96 #Step 6: separating train dataset to 70% for model building and 30% for testing
97 set.seed(1)
98 sample <- sample.int(n = nrow(train), size = floor(.7*nrow(train)), replace = F)
99 train_data <- train[sample, ]
100 train_target <- train[-sample, ]
101
102 #Step 7: Start modeling with all factors
103 #build a model on training data using all variables (the 70%)
104 fit<-lm(count~., train_data)
105 #test model on the training target (30%) to see how good it is
106 predicted.rentals.testing<-predict(fit, train_target)
107 #score: 108.6011
108 rmse(train_target$count,predicted.rentals.testing)
109
110 #let's log(count)
111 logfit<-lm(log(count)~., train_data)
112 #test model on the training target (30%) to see how good it is
113 predicted.rentals.testing<-predict(logfit, train_target)
114 predicted.rentals.testing.nonlog <- exp(predicted.rentals.testing)
115 #score: 0.613963

```



```

116 rmsle(train_target$count,predicted.rentals.testing.nonlog)
117
118 #using Step AIC
119 logfitAIC <- stepAIC(logfit, direction = 'both')
120 #test model on the training target (30%) to see how good it is
121 predicted.rentals.testing<-predict(logfitAIC, train_target)
122 predicted.rentals.testing.nonlog <- exp(predicted.rentals.testing)
123 #score: 0.6140443 - ok this model is worst so we use log(count)~. for now
124 rmsle(train_target$count,predicted.rentals.testing.nonlog)
125
126 #let's use train in caret package with the log(count)~.
127 ctrl <- trainControl(method = "repeatedcv",
128                      number = 3,
129                      repeats = 3,
130                      verboseIter = TRUE,
131                      allowParallel = TRUE)
132 #xgblinear from XGBoost
133 logfit2 <- train(log(count)~., train_data,
134                 method="xgbLinear", trControl = ctrl)
135 #test model on the training target (30%) to see how good it is
136 predicted.rentals.testing<-predict(logfit2, train_target)
137 predicted.rentals.testing.nonlog <- exp(predicted.rentals.testing)
138 #score: 0.422968 - ok this model IS GOODDD
139 rmsle(train_target$count,predicted.rentals.testing.nonlog)
140
141 #LET'S PREDICT
142 predicted.rentals.final<-predict(logfit2, test)
143 predicted.rentals.final.nonlog <- exp(predicted.rentals.final)
144 predicted.rentals = data.frame(datetime = to.predict$datetime, count = predicted.rentals.final.nonlog)
145 colnames(predicted.rentals) <- c("datetime", "count")
146
147 predicted.rentals.final.nonlog <- exp(predicted.rentals.final)
148 predicted.rentals = data.frame(datetime = to.predict$datetime, count = predicted.rentals.final.nonlog)
149 colnames(predicted.rentals) <- c("datetime", "count")
150 write.csv(predicted.rentals, "predicted_rentals.csv", row.names=FALSE)
151 #score on Kaggle 0.52418 [position 1902]
152
153 #attempt 2
154 #let's see restart and see if there is outliers
155 fit<-lm(count~., train)
156 plot(density(resid(fit)))
157 sample_size <- nrow(train)
158 cooks_d <- cooks.distance(fit)
159 #plot cook's distance
160 plot(cooks_d, pch="*", cex=2, main="Influential Obs by Cooks distance")
161 #add cutoff line
162 abline(h = 4/sample_size, col="red")
163 #add labels
164 text(x=1:length(cooks_d)+1, y=cooks_d, labels=ifelse(cooks_d>4/sample_size, names(cooks_d),""), col="red")
165
166 #yes there is, remove them
167 outliers.located <- c(9000, 9009, 9010, 9008, 9011, 8334,
168                     8331, 8332, 8335, 8307, 8333, 8312,5662)
169 #remove the outliers and try modelling again
170 train <- train[-outliers.located, ]
171
172 #separating train dataset to 70% for model building and 30% for testing
173 set.seed(1)
174 sample <- sample.int(n = nrow(train), size = floor(.7*nrow(train)), replace = F)
175 train_data <- train[sample, ]
176 train_target <- train[-sample, ]

```

```

173
174 #xgbLinear from XGBoost #new best!
175 logfit2 <- train(log(count)~., train_data,
176                 method="xgbLinear", trControl = ctrl)
177 #test model on the training target (30%) to see how good it is
178 predicted.rentals.testing<-predict(logfit2, train_target)
179 predicted.rentals.testing.nonlog <- exp(predicted.rentals.testing)
180 #score: 0.4237904 - ok this model looks better
181 rmsle(train_target$count,predicted.rentals.testing.nonlog)
182 #LET'S PREDICT
183 predicted.rentals.final<-predict(logfit2, test)
184 predicted.rentals.final.nonlog <- exp(predicted.rentals.final)
185 predicted.rentals = data.frame(datetime = to.predict$datetime, count = predicted.rentals.final.nonlog)
186 colnames(predicted.rentals) <- c("datetime", "count")
187 write.csv(predicted.rentals, "predicted_rentals.csv", row.names=FALSE)
188 #score on Kaggle 0.50459 omgggg it got better!
189
190 #using stepAIC to find a better model
191 fit<-lm(log(count)~., train_data)
192 logfitAIC <- stepAIC(fit, direction = 'both')
193
194 #ok let's train
195 logfit2 <- train(log(count) ~ season + holiday + workingday + weather + temp +
196                 humidity + windspeed + hour + day + month, train_data,
197                 method="xgbLinear", trControl = ctrl)
198 #test model on the training target (30%) to see how good it is
199 predicted.rentals.testing<-predict(logfit2, train_target)
200 predicted.rentals.testing.nonlog <- exp(predicted.rentals.testing)
201 #score: 0.4231934 - ok this model is a little better but not good enough
202 rmsle(train_target$count,predicted.rentals.testing.nonlog)
203
204 #lets's log the season and hour
205 fit<-lm(log(count) ~ log(as.numeric(season)) + holiday + workingday + weather + temp +
206         humidity + windspeed + log(as.numeric(hour)) + day + month, train_data)
207 plot(fit)
208 #log helps, we get a normal curve
209 plot(density(resid(fit)))
210
211 #ok let's train again
212 logfit2 <- train(log(count) ~ log(as.numeric(season)) + holiday + workingday + weather + temp +
213                 humidity + windspeed + log(as.numeric(hour)) + day + month, train_data,
214                 method="xgbLinear", trControl = ctrl)
215 #test model on the training target (30%) to see how good it is
216 predicted.rentals.testing<-predict(logfit2, train_target)
217 predicted.rentals.testing.nonlog <- exp(predicted.rentals.testing)
218 #score: 0.3432964 - ok this model IS GOODDD
219 rmsle(train_target$count,predicted.rentals.testing.nonlog)
220 #LET'S PREDICT
221 predicted.rentals.final<-predict(logfit2, test)
222 predicted.rentals.final.nonlog <- exp(predicted.rentals.final)
223 predicted.rentals = data.frame(datetime = to.predict$datetime, count = predicted.rentals.final.nonlog)
224 colnames(predicted.rentals) <- c("datetime", "count")
225 write.csv(predicted.rentals, "predicted_rentals.csv", row.names=FALSE)
226 #score on Kaggle 0.49247 omggggggggg [position: 1446]

```